

Portfolio Optimization Techniques for Algorithmic Portfolio Management

Mean-Variance Optimization (MVO / Markowitz)

Conceptual overview

- **Core intuition and objective:** Markowitz's mean-variance framework (1952) aims to maximize expected portfolio return for a given level of risk, or equivalently minimize risk for a given expected return. The risk is modeled via the portfolio variance while expected returns provide the reward. Portfolios are constructed by solving a quadratic optimization problem that trades off return and risk using a **risk-aversion parameter**; a larger risk penalty pushes the solution toward minimum variance portfolios ¹.
- **When and why used:** MVO is the baseline model for modern portfolio theory. It is used to create efficient frontiers and to tailor portfolios for clients with different risk tolerances. Because MVO uses return forecasts, it can allocate more capital to high-alpha assets. However, it is sensitive to forecast error and to the condition number of the covariance matrix ². For large asset universes or unstable covariance matrices, risk-based methods or regularization are often preferred.

Mathematical formulation

- **Objective function (maximization form):**

$$\underset{w}{\text{maximize}} \quad w^\top \mu - \lambda w^\top \Sigma w - c^\top |w - w_0|,$$

where:

- w – vector of portfolio weights (sum to 1).
- μ – vector of expected returns.
- Σ – covariance matrix of returns.
- λ – risk-aversion parameter controlling the trade-off between return and variance.
- $c^\top |w - w_0|$ – optional transaction-cost term penalizing trades from the current holdings w_0 ³.

Equivalent forms include minimizing variance subject to a required return, or maximizing the Sharpe ratio.

- **Constraints:** typical constraints include
- **Fully invested:** $\sum_i w_i = 1$ ⁴.
- **Long-only:** $w_i \geq 0$ ⁴.
- **Leverage limits:** $\|w\|_1 \leq L$ or $\|w\|_1 \leq 1 + \text{max leverage}$ (long-short portfolios). Some practitioners impose gross exposure constraints (e.g., $\sum_i |w_i| \leq 1.5$) to avoid extreme weights.

- **Turnover / transaction cost constraints:** terms such as $\gamma \|w - w_0\|_1$ or $\tau \|w - w_0\|_2^2$ penalize deviations from current holdings to control turnover ⁵.
- **Risk limits:** maximum portfolio volatility or Value-at-Risk; sector or factor exposures; weight bounds.
- **Key assumptions:** returns are (approximately) multivariate normal; covariance matrix is positive definite; forecasts and covariance estimates are unbiased. In practice, these assumptions rarely hold exactly—hence the need for robust estimation and regularization.

Inputs required

1. **Expected returns μ :** can be estimated via historical means, factor models (CAPM, Fama–French), analyst forecasts, or Bayesian models like Black–Litterman ⁶. Because errors in μ have a large effect on portfolio weights ⁷, it is common to shrink forecasts toward a neutral or equilibrium value (see enhancements).
2. **Covariance matrix Σ :** estimated from historical returns (sample covariance), rolling windows, or factor models. Regularization methods such as Ledoit–Wolf shrinkage combine the sample covariance with a scaled identity matrix to improve stability. Shrinkage ensures positive-definiteness and reduces estimation error ⁸.
3. **Risk-free rate:** used if the objective is to maximize the **Sharpe ratio**; the Sharpe ratio compares excess portfolio return to its volatility ⁹.

Algorithm / implementation approach

1. **Estimate inputs:** compute sample mean returns and covariance matrix (optionally using shrinkage). For example, using `pandas` one computes daily returns, then the annualized mean and covariance. For robust estimates, use `scikit-learn`'s `LedoitWolf` estimator.
2. **Formulate optimization problem:** choose a risk-aversion coefficient λ or desired target return. Write the quadratic objective and linear constraints.
3. **Solve with quadratic programming (QP):** the problem is convex and can be solved using quadratic programming solvers (`cvxpy`, `scipy.optimize.minimize` with method `SLSQP`). For long-only MVO, solve

```
import numpy as np
from scipy.optimize import minimize

def mvo_objective(w, mu, cov, lam):
    return -w.dot(mu) + lam * w.dot(cov).dot(w)

constraints = [
    {'type': 'eq', 'fun': lambda w: np.sum(w) - 1},          # fully
    {'type': 'ineq', 'fun': lambda w: w}                    # non-negative
]
result = minimize(mvo_objective, x0=np.ones(n)/n,
```

```
args=(mu, cov, lam), constraints=constraints)
w_opt = result.x
```

Libraries such as **PyPortfolioOpt** provide ready-made classes (`EfficientFrontier`) that accept expected returns and covariance and allow additional objectives (L2 regularization, transaction cost terms). 4. **Add regularization**: to mitigate extreme weights, add an L2 penalty $\lambda_2 \|w\|_2^2$ (ridge) or L1 penalty $\lambda_1 \|w\|_1$ (lasso). `PyPortfolioOpt` supports `L2_reg` to shrink weights toward zero. Group L2 or sector constraints can also be imposed. 5. **Backtesting & turnover control**: simulate rolling-window estimation (e.g., 3–5 years of history) and reoptimize monthly or quarterly. Include transaction cost penalties and gross exposure limits to prevent over-trading ⁵.

Practical challenges & fixes

- **Estimation error in μ and Σ** : small errors in expected returns lead to large changes in weights ² ⁷. Solutions:
 - Use shrinkage estimators (Ledoit–Wolf) for the covariance and Bayesian or shrinkage methods for means.
 - Apply robust optimization: treat μ as uncertain within an ellipsoidal set and optimize the worst-case return ¹⁰.
 - Black–Litterman model combines market equilibrium returns with investor views to produce posterior expected returns ⁶.
- Impose norm or weight constraints to regularize the solution (e.g., limit $\|w\|_1$, use L2 penalty).
- **Covariance instability & ill-conditioning**: high condition number of Σ leads to unstable inverses. Risk-based methods or HRP avoid inverting Σ ². Use shrinkage and factor models to reduce dimensionality.
- **Overfitting & extreme weights**: unconstrained MVO can yield very concentrated portfolios. Fixes: constrain weights between bounds, use a **risk-budget** approach, or add L1/L2 regularization and turnover penalties.
- **Transaction costs & turnover**: include transaction cost terms in the optimization objective or impose turnover constraints. High turnover strategies reduce net performance ⁵. The Markowitz–Ledoit–Wolf paper shows that accounting for stock-specific transaction costs during portfolio selection increases net Sharpe ratios ¹¹.

Enhancements / variants

- **Robust optimization**: treat expected returns as uncertain within an ellipsoidal set. The robust formulation maximizes the worst-case return over the uncertainty set ¹⁰, reducing sensitivity to estimation errors. This often produces more stable portfolios.
- **Bayesian / Black–Litterman**: start from market equilibrium returns (implied by market-cap weights) and adjust them with investor views, weighted by confidence. The resulting posterior means are used in MVO, producing more stable weights ⁶.

- **Norm-constrained optimization:** apply gross exposure constraints ($\|w\|_1 \leq L$) or L2 regularization. This trades off diversification and portfolio sparsity and helps manage transaction costs.
- **Minimum-Tracking Error & other objectives:** incorporate an index or factor-tracking constraint; optimize for maximum Sharpe ratio or risk parity instead of expected return; combine with multi-period dynamic programming for long horizons.

Output interpretation

- **Weights:** represent capital allocations to each asset. Large positive weights indicate high conviction; negative weights (in long-short portfolios) indicate short positions.
- **Portfolio metrics:** compute expected portfolio return $w^\top \mu$, volatility $\sqrt{w^\top \Sigma w}$, and risk-adjusted performance (e.g., Sharpe ratio) ⁹. Additional metrics include maximum drawdown (largest peak-to-trough decline), skewness and kurtosis, and turnover.
- **Sensitivity analysis:** examine how weights change with input perturbations and risk aversion. Use simulation to assess out-of-sample performance and drawdown behaviour.

Minimum Variance Portfolio (MVP / Global Minimum Variance)

Conceptual overview

- **Core intuition:** the MVP seeks the portfolio with the **lowest possible variance** without regard to expected returns. This is equivalent to the MVO with $\lambda \rightarrow \infty$ or with no return term. The strategy appeals to investors whose primary objective is risk minimization.
- **When and why used:** used as a benchmark for low-risk investing and as a building block for multi-asset portfolios (e.g., risk parity). It avoids unstable return forecasts and is purely driven by the covariance structure. However, it tends to allocate heavily to low-volatility assets and may have low expected returns.

Mathematical formulation

- **Objective:** minimize portfolio variance

$$\underset{w}{\text{minimize}} \quad w^\top \Sigma w$$

subject to $\sum_i w_i = 1$ and (for long-only) $w_i \geq 0$ ⁴. Without the non-negativity constraint, the closed-form solution is

$$w^{\text{MVP}} = \frac{\Sigma^{-1} \mathbf{1}}{\mathbf{1}^\top \Sigma^{-1} \mathbf{1}},$$

where $\mathbf{1}$ is the vector of ones ¹². With long-only constraints, solve a quadratic program or sequentially exclude assets with negative weights ¹².

- **Constraints:** fully invested; optional weight bounds; turnover penalties; leverage constraints.
- **Key assumptions:** same as MVO; covariance positive definite; returns have finite variance.

Inputs required

- **Covariance matrix Σ** : the only input. Estimation and shrinkage as described above.
- **No expected returns required**, though some variants impose a minimum return constraint.

Algorithm / implementation

1. **Estimate covariance matrix** (sample or shrinkage). If positive-definite, compute its inverse directly for the closed form. For long-only, apply QP or iterative methods.
2. **Closed-form solution**: compute $\Sigma^{-1}\mathbf{1}$ and normalize.
3. **Long-only variant**: solve `min w.T @ cov @ w` subject to `sum(w)==1` and `w>=0` using `scipy.optimize.minimize` or `cvxpy`.
4. **Regularization / sparse MVP**: add small L2 penalty to ensure positive definiteness or L1 penalty to promote sparsity.

Practical challenges & fixes

- **Concentration**: MVP often concentrates in a few low-volatility assets. Add maximum weight constraints or integrate expected returns.
- **Estimation error**: shrinkage and factor models mitigate covariance estimation error.
- **Turnover**: as with MVO, include turnover penalties.

Enhancements / variants

- **Target return MVP**: impose a minimum expected return to avoid extremely low-return portfolios.
- **Risk budgeting / risk parity**: allocate risk contributions rather than minimize variance (see next section).
- **Extended MVP with constraints**: incorporate sector caps, ESG constraints, or leverage limits.

Output interpretation

- **Weights** often overweight low-volatility assets (e.g., bonds). Examine risk contribution of each asset.
- **Portfolio metrics**: extremely low volatility; measure Sharpe ratio to ensure return commensurate with risk.

Risk Parity (Equal Risk Contribution)

Conceptual overview

- **Core intuition**: risk parity aims to allocate capital so that each asset contributes **equal proportion of risk** to the portfolio. Risk contribution of asset i is defined as

$$\text{RRC}_i = \frac{w_i(\Sigma w)_i}{w^\top \Sigma w},$$

and risk parity requires $\text{RRC}_i = 1/N$ for N assets ¹³. Unlike MVO, it avoids estimating expected returns and reduces concentration by balancing risk exposure ¹⁴.

- **When and why used:** risk parity is popular among hedge funds (e.g., Bridgewater's All Weather fund). It performs well during uncertain markets and aims for steady growth by allocating more capital to low-volatility assets ¹⁵. Because it does not rely on return forecasts, it is less sensitive to estimation error and provides good diversification.

Mathematical formulation

- **Objective:** find weights w that minimize the variance of the risk contributions—i.e., minimize

$$f(w) = \sum_{i=1}^N \left(\text{RRC}_i - \frac{1}{N} \right)^2.$$

Equivalent formulation is to minimize the sum of squared differences of risk contributions ¹³. Long-only and fully invested constraints are usually imposed.

- **Constraints:**
 - $\sum_i w_i = 1$
 - $w_i \geq 0$ (long-only) or with leverage.
 - Optionally risk budgets b_i instead of equal risk: $\text{RRC}_i = b_i$ with $\sum b_i = 1$.
- **Key assumptions:** covariance matrix exists but need not be invertible; risk is measured via variance. Extensions use VaR, CVaR or other coherent risk measures.

Inputs required

- **Covariance matrix Σ :** only input. Estimation and shrinkage as before.

Algorithm / implementation approach

1. **Compute covariance matrix** (sample or shrinkage). Standardize for annual frequency.
2. **Define risk contribution and objective:** implement functions to compute portfolio variance, marginal risk contributions, and risk contributions. Example from the LuxAlgo article ¹⁶:

```
import numpy as np
from scipy.optimize import minimize

def portfolio_variance(w, cov):
    return w.T @ cov @ w

def risk_contribution(w, cov):
    var_p = portfolio_variance(w, cov)
    marg = cov @ w
```

```

    return (w * marg) / var_p

def risk_parity_objective(w, cov):
    rc = risk_contribution(w, cov)
    return np.sum((rc - rc.mean()) ** 2)

constraints = [
    {'type': 'eq', 'fun': lambda w: w.sum() - 1},
    {'type': 'ineq', 'fun': lambda w: w}
]
init_w = np.ones(N)/N
res = minimize(risk_parity_objective, init_w, args=(cov,), method='SLSQP',
               constraints=constraints, options={'ftol': 1e-12})
rp_weights = res.x

```

3. **Implement risk budgets:** modify the objective to target custom risk budgets b_i . Minimize $\sum_i (\text{RRC}_i - b_i)^2$ subject to $\sum b_i = 1$.
4. **Hierarchical Risk Parity (HRP):** a machine-learning algorithm that builds risk parity portfolios without inverting Σ . HRP addresses the problem of numerical instability (Markowitz's curse) by:
5. **Hierarchical clustering:** cluster assets using distances derived from correlations ¹⁷. Compute pairwise distances $d_{i,j} = \sqrt{\frac{1}{2}(1 - \rho_{ij})}$ and use linkage clustering ¹⁷.
6. **Quasi-diagonalization:** reorder the correlation matrix according to the clustering to expose block structure ¹⁷.
7. **Recursive bisection:** allocate weights recursively by splitting the portfolio into clusters and distributing capital in proportion to cluster variances; within each cluster, allocate weights inversely proportional to asset variances ¹⁷.

HRP avoids matrix inversion, is robust to estimation error and tends to deliver lower out-of-sample variance ¹⁷. 5. **Implementation libraries:** Python libraries such as `riskfolio-lib` implement both classical risk parity and HRP. HRP can also be coded from scratch using `scipy.cluster.hierarchy` for clustering.

Practical challenges & fixes

- **Choice of risk measure:** variance may not capture tail risk. CVaR risk parity extends risk contributions to CVaR; however, research suggests that equalizing tail-risk contributions does not always improve performance ¹⁴.
- **Covariance estimation:** use robust estimators (shrinkage, factor models); HRP less sensitive to ill-conditioned Σ .
- **Concentration:** naive risk parity (inverse volatility weighting) may allocate heavily to low-volatility assets (e.g., bonds) ¹⁴. Use risk budgets or minimum weight constraints to prevent extreme allocations.

Enhancements / variants

- **Group risk parity:** equalize risk contributions among groups (e.g., asset classes) and then within groups. Useful in multi-asset portfolios.
- **Factor risk parity:** equalize risk contributions to systematic factors.
- **Tail risk parity:** define risk contributions using CVaR or expected shortfall. Implementation is more complex and requires simulation.
- **Dynamic risk parity:** adjust risk budgets based on market regimes (e.g., increase equity risk during low volatility periods).
- **Hierarchical Risk Parity (HRP)** – described above – addresses estimation error and has become popular; Monte Carlo simulations indicate HRP achieves lower out-of-sample variance than the critical line algorithm ¹⁸.

Output interpretation

- **Weights:** risk parity weights often overweight low-volatility assets relative to market cap. Evaluate risk contributions to verify equalization.
- **Portfolio metrics:** compute portfolio volatility, risk contribution per asset, Sharpe ratio, and maximum drawdown. Compare to equal-weight and MVO portfolios.

Equal-Weight Portfolio

Conceptual overview

- **Core intuition:** allocate equal capital to each asset; no estimation of returns or covariances is required. This naive diversification often performs surprisingly well due to diversification benefits and robustness to estimation error ¹⁹.
- **When and why used:** used as a benchmark or starting allocation. It provides maximum diversification and inherently rebalances by selling winners and buying losers, giving it a value tilt ²⁰.

Mathematical formulation

- **Weights:** $w_i = 1/N$ for all assets ¹⁹.
- **Constraints:** fully invested; may include long-only; no other constraints are needed.
- **No optimization required.**

Inputs required

- None beyond list of assets. Because no parameter estimation is required, the equal-weight portfolio is immune to estimation error.

Implementation

- Simply assign each asset weight $1/N$ and rebalance periodically to maintain equal weights. Example in Python:


```
N = len(assets)
eq_weights = np.ones(N) / N
```

- **Rebalancing:** due to price movements, weights drift; rebalance monthly, quarterly or when weights deviate by more than a tolerance.

Practical challenges & fixes

- **High turnover:** equal weighting requires frequent rebalancing, leading to transaction costs ²¹. To reduce turnover, use tolerance bands (rebalance only if a weight drifts beyond $\pm$ threshold). Time-based rebalancing strategies (quarterly or annually) are easy to implement but may incur unnecessary trades.
- **Concentration risk:** equal weight overweights small-cap or illiquid assets relative to their market caps. Set liquidity filters or weight caps to avoid illiquid exposures.

Enhancements / variants

- **Value-weighted or risk-adjusted equal weight:** adjust weights by fundamentals or volatility; an **inverse volatility** equal weight is a naive risk parity approach.
- **Custom weighting schemes:** equal weight within sectors or factors and allocate across groups differently.

Output interpretation

- **Weights:** simple 1/N allocation. Portfolio metrics: compute average return, volatility, Sharpe ratio and drawdown. Equal weight often exhibits higher volatility but can outperform cap-weighted indices after adjusting for risk and transaction costs ²¹.

Comparison Across Methods

Method	Strengths	Weaknesses	Suitable Market Regimes
Mean-Variance Optimization	Incorporates expected returns and provides flexible risk-return trade-off; can tailor to investor views; supports constraints and regularization	Highly sensitive to estimation errors in expected returns and covariance; may produce extreme weights; requires stable input estimates ² 7	Works well when return forecasts are accurate and stable; risk-averse investors willing to accept higher turnover and modelling complexity
Minimum Variance Portfolio	Requires only covariance matrix; avoids return forecasts; closed-form solution; reduces volatility	May concentrate in low-volatility assets and deliver low expected returns; still sensitive to covariance estimation error	Suitable in volatile or uncertain markets where capital preservation is key; baseline low-risk benchmark

Method	Strengths	Weaknesses	Suitable Market Regimes
Risk Parity	Balances risk contributions across assets; does not require expected returns; robust to estimation error; HRP variant avoids matrix inversion ¹⁷	May heavily allocate to low-volatility assets (bonds), underweighting high-risk assets; objective is implicit and may not maximize return; requires iterative optimization	Works well in diversified multi-asset portfolios, especially during economic downturns or when covariance matrices are ill-conditioned
Equal Weight	Simplest and most robust; no parameter estimation; provides diversification and value tilt ²⁰	Ignores risk differences; high turnover to maintain equal weights; overweight of small or illiquid assets	Suitable as a benchmark or core allocation; performs well when forecasting is difficult or when assets are similarly risky

Production Considerations

Rebalancing frequency and turnover control

- **Time-based vs. threshold-based rebalancing:** time-based rebalancing (e.g., quarterly or annually) reviews the portfolio at fixed intervals. It is simple but may generate unnecessary trades. Threshold (percentage-of-portfolio) rebalancing sets tolerance bands (e.g., $\pm 5\%$ deviation) and rebalances only when weights drift outside the bands. This approach reduces turnover and transaction costs.
- **Consider transaction costs and taxes:** high rebalancing frequency increases trading fees and potential tax liabilities. Model transaction costs explicitly in the optimization objective (e.g., $c_i \|w - w_0\|_1$ penalties) and calibrate penalty parameters based on desired turnover ⁵. Long-only portfolios with heavy turnover may see lower net performance ⁵.
- **Dynamic vs. constant rebalancing:** dynamic rebalancing responds to market conditions (e.g., volatility regime changes) or uses adaptive risk budgets. Constant proportion portfolio insurance (CPPI) adjusts allocations based on portfolio value relative to a floor.
- **Backtesting and evaluation:** use rolling windows for estimation (e.g., five years of data) and evaluation periods to select penalty parameters. A rolling or nested cross-validation approach can calibrate turnover penalties ⁵.

Handling transaction costs

- **Incorporate transaction costs in optimization:** directly include stock-specific bid-ask spreads or fixed trading costs in the objective. The Markowitz portfolios under transaction costs paper shows that accounting for transaction costs at the portfolio selection stage leads to higher net Sharpe ratios and reduces turnover ¹¹. Many academic studies neglect this and therefore overestimate performance ⁵.
- **Turnover penalties:** add L1 or L2 penalties on weight changes $\|w - w_0\|$ to control trading activity. A penalty parameter allows the manager to choose a comfortable turnover level and can be calibrated via backtesting ⁵.

Handling missing data and illiquid assets

- **Missing returns:** use algorithms that can estimate means and covariances with missing data instead of discarding observations. The MATLAB example demonstrates two approaches: `nanskip` (drop rows with any missing values) and expectation conditional maximization (ECM) to estimate means and covariance using all available data. The ECM approach converges even with large gaps, such as 87% missing data for Google.
- **Synchronize data:** when assets trade in different time zones or have different trading days, align them to a common calendar by forward filling or interpolation ²².
- **Illiquid assets:** private assets can have long lock-up periods. Illiquidity prevents continuous rebalancing and introduces uncertainty in portfolio allocations. The SIERF policy brief highlights that ignoring illiquidity risk can lead to substantial welfare losses; the inability to rebalance during lock-ups causes unintended overallocation to private assets ²³. A dynamic model must incorporate expected waiting times for liquidity events and adjust strategic allocations accordingly ²⁴. Investors may reduce target weights for severely illiquid assets ²⁵.

Implementation and scaling considerations

- **Scalability:** for large universes (hundreds of assets), use efficient linear algebra (NumPy, pandas) and specialized solvers (e.g., `cvxpy` with `OSQP` or `MOSEK`). HRP scales to large dimensions without inverting Σ .
- **Parallel computation:** compute sample means/covariances and optimize portfolios in parallel across rolling windows.
- **Backtesting and simulation:** implement vectorized backtests with realistic slippage, transaction costs, and rebalancing schedules. Evaluate metrics like return, volatility, Sharpe ratio, drawdown and turnover.
- **Integration with data pipelines:** integrate data retrieval (e.g., via `yfinance` or database) and handle missing data (forward fill/interpolation). Use stable identifiers (tickers) and handle corporate actions (splits, dividends).

Summary and recommendations

- **Choose the method based on objectives:** for return-seeking strategies where forecasts are available, MVO with robust estimation and regularization may offer the best trade-off. For risk minimization or when forecasts are unreliable, MVP or risk parity provide more stable allocations. Equal-weight is a robust baseline and often competitive after costs.
- **Mitigate estimation error:** shrinkage and Bayesian techniques for returns and covariance; robust optimization; HRP to avoid matrix inversion.
- **Control turnover and costs:** incorporate transaction cost penalties; calibrate penalties to achieve desired turnover levels; use threshold-based rebalancing and consider tax implications.
- **Handle data issues:** use algorithms capable of dealing with missing data; adjust for illiquidity; ensure data quality and synchronization.
- **Evaluate across regimes:** test strategies in bull and bear markets, high and low volatility regimes, and under different interest-rate environments. Use out-of-sample testing and cross-validation to avoid overfitting.

By combining these techniques with thoughtful estimation and practical constraints, a portfolio manager can build an algorithmic engine that dynamically optimizes portfolios for a wide range of market conditions and investment objectives.

1 Mean-Variance Optimization | QuestDB

<https://questdb.com/glossary/mean-variance-optimization/>

2 17 18 Hierarchical Risk Parity - Wikipedia

https://en.wikipedia.org/wiki/Hierarchical_Risk_Parity

3 markowitz.pdf

<https://web.stanford.edu/~boyd/papers/pdf/markowitz.pdf>

4 12 Understanding the Long-Only Minimum Variance Portfolio

<https://arxiv.org/pdf/2603.07692>

5 11 Markowitz portfolios under transaction costs

https://www.econ.uzh.ch/dam/jcr:cd9c657f-16d3-4155-872d-e5ebae7557ac/qref_2025.pdf

6 Portfolio Optimization — Black-Litterman, AI & Advanced Techniques Guide

<https://marketxls.com/blog/advanced-portfolio-optimization-black-litterman-ai>

7 10 Portfolio_Optimization_with_Estimation_Errors.pdf

https://optimization-online.org/wp-content/uploads/2024/12/Portfolio_Optimization_with_Estimation_Errors.pdf

8 skfolio.moments.LedoitWolf - skfolio

<https://skfolio.org/generated/skfolio.moments.LedoitWolf.html>

9 Sharpe Ratio: Meaning, Formula, Advantages and Limitations | m.Stock

<https://www.mstock.com/articles/what-is-sharpe-ratio>

13 Portfolio Optimization - Risk Parity Portfolios

<https://portfoliooptimizationbook.com/slides/slides-rpp.pdf>

14 Risk Parity Asset Allocation - QuantPedia

<https://quantpedia.com/risk-parity-asset-allocation/>

15 16 22 Risk Parity Allocation with Python

<https://www.luxalgo.com/blog/risk-parity-allocation-with-python/>

19 ErnstThompsonMiao.pdf

<https://www.ma.imperial.ac.uk/~pernst/ErnstThompsonMiao.pdf>

20 21 Equal-Weighted Index - Definition, Pros/Cons

<https://corporatefinanceinstitute.com/resources/equities/equal-weighted-index/>

23 24 25 SUERF-Policy-Brief-1122_Dimitrov.pdf

https://www.suerf.org/wp-content/uploads/2025/03/SUERF-Policy-Brief-1122_Dimitrov.pdf